

Classification With Transfer Learning

1 Importing Necessary Libraries

```
[1]: import os
import numpy as np
import seaborn as sns
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.utils import plot_model
from tensorflow.keras.applications import VGG19
from tensorflow.keras import layers, mixed_precision
from sklearn.metrics import classification_report, confusion_matrix
```

2 Defining Model Architecture (VGG19) and Data Loading

```
[4]: def cnn(input_shape, num_classes):

    base = VGG19(weights = "imagenet", include_top = False, input_shape =
↪input_shape)
    base.trainable = False

    model = tf.keras.Sequential([
        base,
        # layers.GlobalAveragePooling2D(),
        layers.Flatten(),

        layers.Dense(512, activation = "relu"),
        layers.BatchNormalization(),

        layers.Dense(256, activation = "relu"),
        layers.BatchNormalization(),

        layers.Dense(128, activation = "relu"),
        layers.BatchNormalization(),

        layers.Dense(num_classes, activation = "softmax")
    ])
```

```
return model
```

```
[6]: dir = "../Dataset/Intel/seg_train"

batch_size = 16
img_h, img_w = 150, 150

train_dataset = tf.keras.utils.image_dataset_from_directory(dir, seed = 100,
↪image_size = (img_h, img_w), batch_size = batch_size, shuffle = True)
```

Found 14034 files belonging to 6 classes.

```
[7]: train_dataset.class_names
```

```
[7]: ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

3 Model Initialization

```
[14]: model = cnn(input_shape = (150, 150, 3), num_classes = 6)
```

3.1 Model Architecture Summary and Visualization

```
[27]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 4, 4, 512)	20024384
flatten_2 (Flatten)	(None, 8192)	0
dense_8 (Dense)	(None, 512)	4194816
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dense_9 (Dense)	(None, 256)	131328
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dense_10 (Dense)	(None, 128)	32896
batch_normalization_8 (Batch Normalization)	(None, 128)	512

```
dense_11 (Dense)                (None, 6)                774
```

```
=====
```

```
Total params: 24,387,782
```

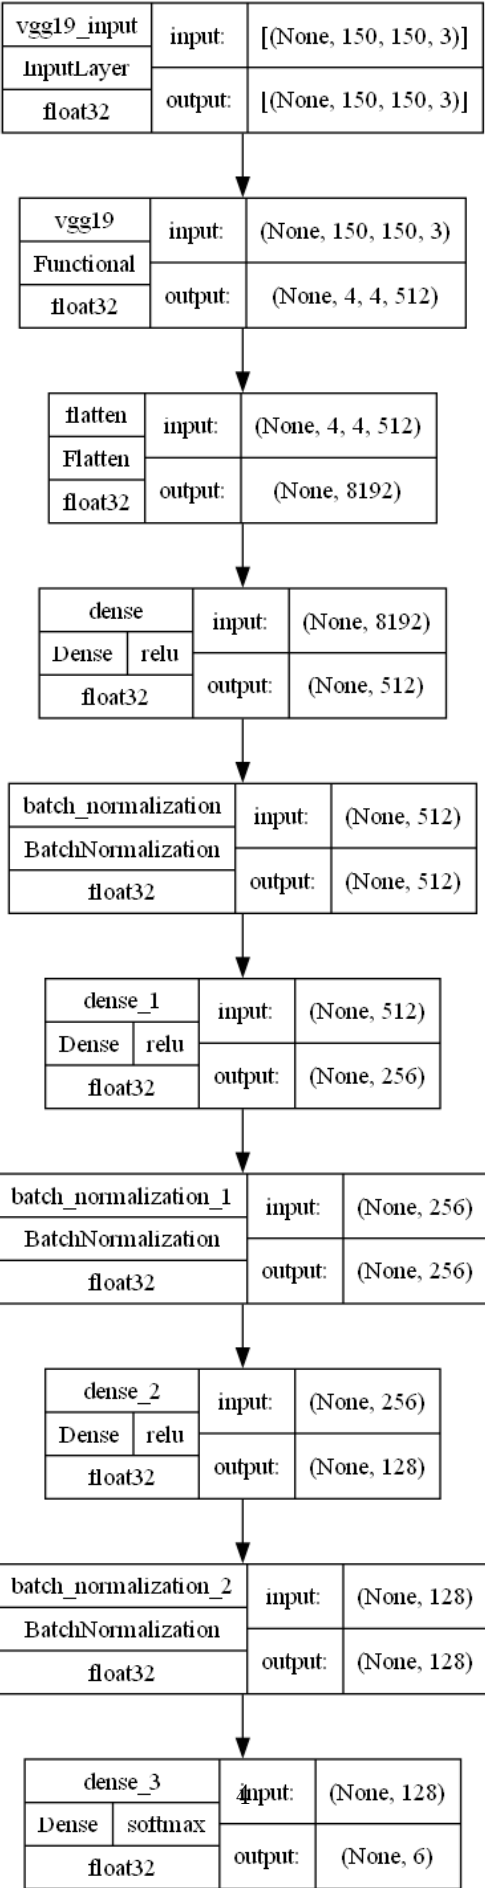
```
Trainable params: 4,361,606
```

```
Non-trainable params: 20,026,176
```

```
-----
```

```
[16]: plot_model(model, show_layer_names = True, show_dtype = True,
      ↪ show_layer_activations = True, show_shapes = True)
```

```
[16]:
```



3.2 Training

```
[29]: model.compile(optimizer = "adam", loss = "sparse_categorical_crossentropy",  
    ↪ metrics = ["accuracy"])
```

```
[31]: # Ensure GPU Usage (if you want to use CPU, "/CPU:0")  
with tf.device("/GPU:0"):  
    history = model.fit(train_dataset, epochs = 20)
```

```
Epoch 1/20  
878/878 [=====] - 116s 125ms/step - loss: 0.4186 -  
accuracy: 0.8511  
Epoch 2/20  
878/878 [=====] - 94s 107ms/step - loss: 0.2434 -  
accuracy: 0.9139  
Epoch 3/20  
878/878 [=====] - 98s 111ms/step - loss: 0.1806 -  
accuracy: 0.9344  
Epoch 4/20  
878/878 [=====] - 94s 107ms/step - loss: 0.1376 -  
accuracy: 0.9510  
Epoch 5/20  
878/878 [=====] - 67s 77ms/step - loss: 0.1153 -  
accuracy: 0.9577  
Epoch 6/20  
878/878 [=====] - 68s 77ms/step - loss: 0.0973 -  
accuracy: 0.9642  
Epoch 7/20  
878/878 [=====] - 68s 77ms/step - loss: 0.0766 -  
accuracy: 0.9720  
Epoch 8/20  
878/878 [=====] - 68s 78ms/step - loss: 0.0631 -  
accuracy: 0.9785  
Epoch 9/20  
878/878 [=====] - 68s 78ms/step - loss: 0.0555 -  
accuracy: 0.9803  
Epoch 10/20  
878/878 [=====] - 70s 79ms/step - loss: 0.0540 -  
accuracy: 0.9808  
Epoch 11/20  
878/878 [=====] - 71s 81ms/step - loss: 0.0419 -  
accuracy: 0.9849  
Epoch 12/20  
878/878 [=====] - 40s 46ms/step - loss: 0.0417 -  
accuracy: 0.9866
```

```

Epoch 13/20
878/878 [=====] - 41s 47ms/step - loss: 0.0330 -
accuracy: 0.9889
Epoch 14/20
878/878 [=====] - 42s 48ms/step - loss: 0.0293 -
accuracy: 0.9900
Epoch 15/20
878/878 [=====] - 42s 48ms/step - loss: 0.0329 -
accuracy: 0.9885
Epoch 16/20
878/878 [=====] - 42s 48ms/step - loss: 0.0292 -
accuracy: 0.9905
Epoch 17/20
878/878 [=====] - 42s 48ms/step - loss: 0.0272 -
accuracy: 0.9912
Epoch 18/20
878/878 [=====] - 43s 48ms/step - loss: 0.0254 -
accuracy: 0.9911
Epoch 19/20
878/878 [=====] - 43s 49ms/step - loss: 0.0253 -
accuracy: 0.9917
Epoch 20/20
878/878 [=====] - 43s 49ms/step - loss: 0.0203 -
accuracy: 0.9932

```

4 Saving The Trained Model

```
[179]: model.save("../Models/Intel/Model With Transfer Learning.h5")
```

5 Model Evaluation and Other Metrics

5.1 Loss and Accuracy Graphs

```

[108]: # Plot training accuracy and loss
def plot_loss_accuracy(history):
    acc = history.history['accuracy']
    loss = history.history['loss']

    epochs = range(1, len(acc) + 1)

    plt.figure(figsize=(10, 5)) # Adjust size

    # Plot accuracy and loss on the same graph
    plt.plot(epochs, acc, label='Training Accuracy') # Blue line with circles
    plt.plot(epochs, loss, label='Training Loss') # Red line with triangles

```

```


# Labels & title
plt.xlabel('Epochs')
plt.ylabel('Metrics')
plt.title('Training Accuracy & Loss')

# Grid and legend
plt.grid(True, linestyle='--', alpha=0.6) # Dashed gridlines
plt.legend()

plt.show()

plot_loss_accuracy(history)

```



5.2 Confusion Matrix and Classification Report

```

[100]: dir = "../Dataset/Intel/seg_test"

batch_size = 16
img_h, img_w = 150, 150

test_dataset = tf.keras.utils.image_dataset_from_directory(dir, seed = 100,
↪ image_size = (img_h, img_w), batch_size = batch_size, shuffle = False)

```

Found 3000 files belonging to 6 classes.

```

[110]: y_true = np.concatenate([labels.numpy() for _, labels in test_dataset])

```

```

# Get model predictions
y_pred = model.predict(test_dataset)
y_pred_classes = np.argmax(y_pred, axis=1)

# Define class names
class_names = test_dataset.class_names

# Generate Classification Report
print(" Classification Report:\n")
print(classification_report(y_true, y_pred_classes, target_names=class_names))

# Generate Confusion Matrix
cm = confusion_matrix(y_true, y_pred_classes)

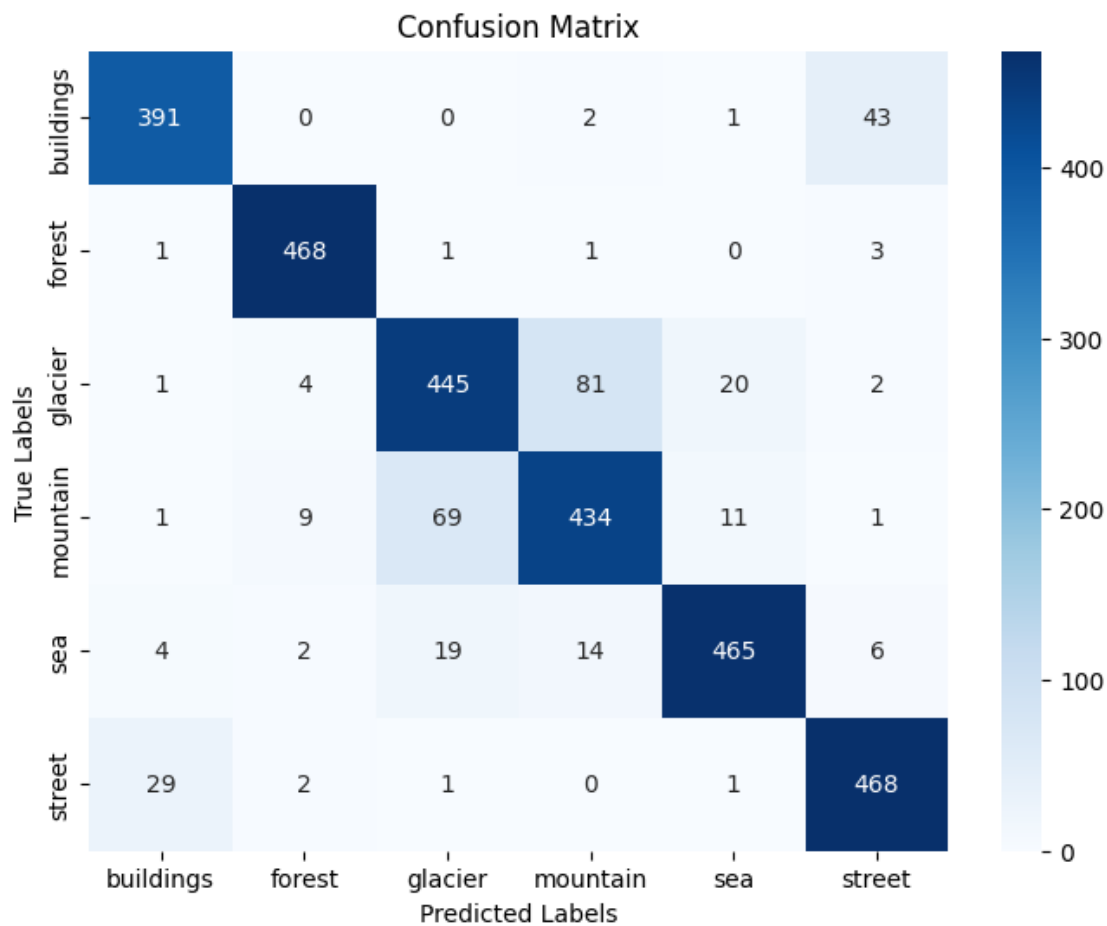
# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names,
            yticklabels=class_names)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```

188/188 [=====] - 8s 41ms/step

Classification Report:

	precision	recall	f1-score	support
buildings	0.92	0.89	0.91	437
forest	0.96	0.99	0.98	474
glacier	0.83	0.80	0.82	553
mountain	0.82	0.83	0.82	525
sea	0.93	0.91	0.92	510
street	0.89	0.93	0.91	501
accuracy			0.89	3000
macro avg	0.89	0.89	0.89	3000
weighted avg	0.89	0.89	0.89	3000



6 Inference

6.1 Importing the necessary libraries

```
[135]: import cv2
import numpy as np
from tensorflow.keras.models import load_model
import matplotlib.pyplot as plt
import random
import math
```

```
[113]: model = load_model("../Models/Intel/Model With Transfer Learning.h5")
```

```
[149]: def inference_grid(img_paths, grid_cols=3):
    class_names = ['buildings', 'forest', 'glacier', 'mountain', 'sea',
↪ 'street']
```

```

num_images = len(img_paths)
grid_rows = math.ceil(num_images / grid_cols) # Calculate number of rows
↳needed

plt.figure(figsize=(grid_cols * 4, grid_rows * 4)) # Adjust figure size
↳dynamically

for i, img_path in enumerate(img_paths):
    img = cv2.imread(img_path)

    # Preprocess image
    img_expanded = np.expand_dims(img, axis=0)

    prediction = model.predict(img_expanded)
    predicted_class = np.argmax(prediction, axis=1)[0]
    predicted_label = class_names[predicted_class]

    # Plot image in a grid
    plt.subplot(grid_rows, grid_cols, i + 1)
    plt.imshow(img)
    plt.axis('off')
    plt.title(f"{predicted_label}", fontsize=12)

plt.tight_layout() # Adjust layout for better spacing
plt.show()

```

```

[174]: parent_dir = "../Dataset/Intel/seg_pred/"

num_images = 9 # Number of images to predict

img_files = random.sample(os.listdir(parent_dir), num_images)
img_paths = [os.path.join(parent_dir, img) for img in img_files]

inference_grid(img_paths, grid_cols = 3)

```

```

1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step

```

buildings



glacier




street




sea



sea



glacier



forest



street



buildings

