

Classification without Transfer Learning

1 Importing Necessary Libraries

```
[23]: import os
import numpy as np
import seaborn as sns
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.utils import plot_model
from tensorflow.keras.applications import VGG19
from tensorflow.keras import layers, mixed_precision
from sklearn.metrics import classification_report, confusion_matrix
```

2 Defining Model Architecture and Data Loading

```
[9]: def cnn(input_shape, num_classes):
    model = tf.keras.Sequential([
        layers.Conv2D(64, (3, 3), activation = "relu", padding = "same",
        ↪input_shape = input_shape),
        layers.BatchNormalization(),
        layers.Conv2D(64, (3, 3), activation = "relu", padding = "same"),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        layers.Conv2D(128, (3, 3), activation = "relu", padding = "same"),
        layers.BatchNormalization(),
        layers.Conv2D(128, (3, 3), activation = "relu", padding = "same"),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),

        # Can include these layers for more complexity
        # layers.Conv2D(256, (3, 3), activation = "relu", padding = "same"),
        # layers.BatchNormalization(),
        # layers.Conv2D(256, (3, 3), activation = "relu", padding = "same"),
        # layers.BatchNormalization(),
        # layers.MaxPooling2D((2, 2)),
```

```

# layers.Conv2D(512, (3, 3), activation = "relu", padding = "same"),
# layers.BatchNormalization(),
# layers.Conv2D(512, (3, 3), activation = "relu", padding = "same"),
# layers.BatchNormalization(),
# layers.MaxPooling2D((2, 2)),

# layers.Conv2D(1024, (3, 3), activation = "relu", padding = "same"),
# layers.BatchNormalization(),
# layers.Conv2D(1024, (3, 3), activation = "relu", padding = "same"),
# layers.BatchNormalization(),
# layers.MaxPooling2D((2, 2)),

layers.Flatten(),
layers.Dense(128, activation = "relu"),
layers.BatchNormalization(),

layers.Dense(64, activation = "relu"),
layers.BatchNormalization(),

layers.Dense(32, activation = "relu"),
layers.BatchNormalization(),

layers.Dense(num_classes, activation = "softmax")
])

return model

```

```

[12]: dir = "../Dataset/Intel/seg_train/"

batch_size = 16
img_h, img_w = 150, 150

train_dataset = tf.keras.utils.image_dataset_from_directory(dir, seed = 100,
↳ image_size = (img_h, img_w), batch_size = batch_size, shuffle = True)

```

Found 14034 files belonging to 6 classes.

3 Model Initialization

```

[14]: model = cnn(input_shape = (150, 150, 3), num_classes = 6)

```

3.1 Model Architecture Summary and Visualization

```

[16]: model.summary()

```

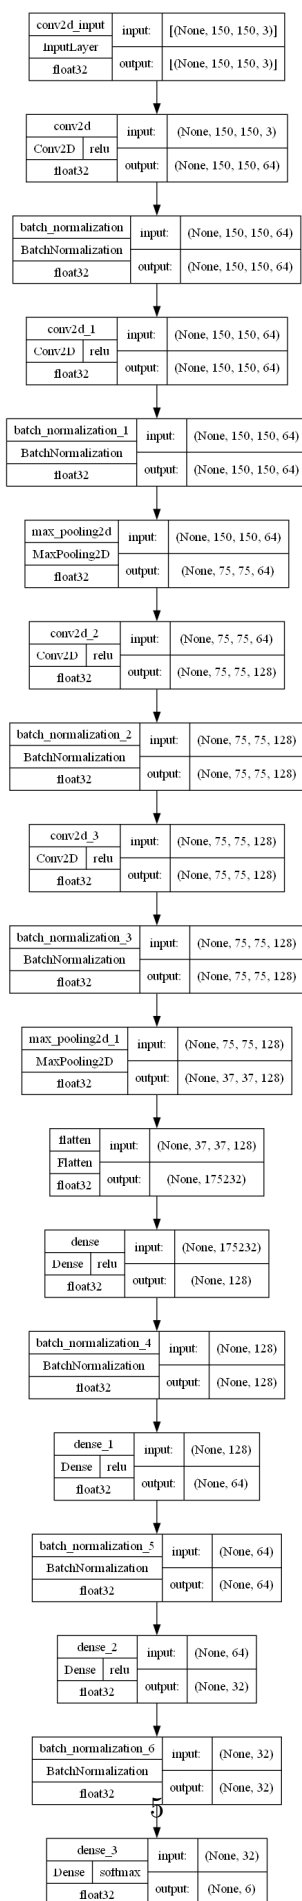
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 64)	1792
batch_normalization (Batch Normalization)	(None, 150, 150, 64)	256
conv2d_1 (Conv2D)	(None, 150, 150, 64)	36928
batch_normalization_1 (Batch Normalization)	(None, 150, 150, 64)	256
max_pooling2d (MaxPooling2D)	(None, 75, 75, 64)	0
conv2d_2 (Conv2D)	(None, 75, 75, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 75, 75, 128)	512
conv2d_3 (Conv2D)	(None, 75, 75, 128)	147584
batch_normalization_3 (Batch Normalization)	(None, 75, 75, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 128)	0
flatten (Flatten)	(None, 175232)	0
dense (Dense)	(None, 128)	22429824
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dense_1 (Dense)	(None, 64)	8256
batch_normalization_5 (Batch Normalization)	(None, 64)	256
dense_2 (Dense)	(None, 32)	2080
batch_normalization_6 (Batch Normalization)	(None, 32)	128
dense_3 (Dense)	(None, 6)	198

Total params: 22,702,950
Trainable params: 22,701,734
Non-trainable params: 1,216

```
[25]: plot_model(model, show_layer_names = True, show_dtype = True, ↵  
      ↪ show_layer_activations = True, show_shapes = True)
```

[25]:



3.2 Training

```
[15]: model.compile(optimizer = "adam", loss = "sparse_categorical_crossentropy",  
    ↪ metrics = ["accuracy"])
```

```
[17]: with tf.device("/GPU:0"):  
    history = model.fit(train_dataset, epochs = 10)
```

```
Epoch 1/10  
878/878 [=====] - 88s 96ms/step - loss: 0.8954 -  
accuracy: 0.6638  
Epoch 2/10  
878/878 [=====] - 87s 99ms/step - loss: 0.6277 -  
accuracy: 0.7701  
Epoch 3/10  
878/878 [=====] - 87s 99ms/step - loss: 0.5707 -  
accuracy: 0.7959  
Epoch 4/10  
878/878 [=====] - 88s 100ms/step - loss: 0.4522 -  
accuracy: 0.8353  
Epoch 5/10  
878/878 [=====] - 88s 100ms/step - loss: 0.3931 -  
accuracy: 0.8603  
Epoch 6/10  
878/878 [=====] - 88s 100ms/step - loss: 0.3321 -  
accuracy: 0.8826  
Epoch 7/10  
878/878 [=====] - 88s 100ms/step - loss: 0.2473 -  
accuracy: 0.9147  
Epoch 8/10  
878/878 [=====] - 88s 100ms/step - loss: 0.1901 -  
accuracy: 0.9329  
Epoch 9/10  
878/878 [=====] - 87s 99ms/step - loss: 0.1389 -  
accuracy: 0.9528  
Epoch 10/10  
878/878 [=====] - 87s 99ms/step - loss: 0.1204 -  
accuracy: 0.9605
```

4 Saving The Model

```
[62]: model.save("../Models/Intel/Model Without Transfer Learning.h5")
```

5 Model Evaluation and Other Metrics

5.1 Loss and Accuracy Graphs

```
[66]: # Plot training accuracy and loss
def plot_loss_accuracy(history):
    acc = history.history['accuracy']
    loss = history.history['loss']

    epochs = range(1, len(acc) + 1)

    plt.figure(figsize=(10, 5)) # Adjust size

    # Plot accuracy and loss on the same graph
    plt.plot(epochs, acc, label='Training Accuracy') # Blue line with circles
    plt.plot(epochs, loss, label='Training Loss') # Red line with triangles

    # Labels & title
    plt.xlabel('Epochs')
    plt.ylabel('Metrics')
    plt.title('Training Accuracy & Loss')

    # Grid and legend
    plt.grid(True, linestyle='--', alpha=0.6) # Dashed gridlines
    plt.legend()

    plt.show()

plot_loss_accuracy(history)
```



5.2 Confusion Matrix and Classification Report

```
[69]: dir = "../Dataset/Intel/seg_test"

batch_size = 16
img_h, img_w = 150, 150

test_dataset = tf.keras.utils.image_dataset_from_directory(dir, seed = 100,
    ↳ image_size = (img_h, img_w), batch_size = batch_size, shuffle = False)
```

Found 3000 files belonging to 6 classes.

```
[75]: y_true = np.concatenate([labels.numpy() for _, labels in test_dataset])

# Get model predictions
y_pred = model.predict(test_dataset)
y_pred_classes = np.argmax(y_pred, axis=1)

# Define class names
class_names = test_dataset.class_names

# Generate Classification Report
print(" Classification Report:\n")
print(classification_report(y_true, y_pred_classes, target_names=class_names))

# Generate Confusion Matrix
cm = confusion_matrix(y_true, y_pred_classes)


# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names,
    ↳ yticklabels=class_names)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```

188/188 [=====] - 2s 13ms/step

Classification Report:

	precision	recall	f1-score	support
buildings	0.86	0.80	0.83	437
forest	0.93	0.96	0.94	474
glacier	0.79	0.66	0.72	553

mountain	0.78	0.75	0.77	525
sea	0.75	0.88	0.81	510
street	0.83	0.88	0.86	501
accuracy			0.82	3000
macro avg	0.82	0.82	0.82	3000
weighted avg	0.82	0.82	0.82	3000



6 Inference

6.1 Importing the necessary libraries

```
[52]: import cv2
import numpy as np
from tensorflow.keras.models import load_model
import matplotlib.pyplot as plt
import random
```

```
import math
```

```
[54]: model = load_model("../Models/Intel/Model Without Transfer Learning.h5")
```

```
INFO:tensorflow:Mixed precision compatibility check (mixed_float16): OK  
Your GPU will likely run quickly with dtype policy mixed_float16 as it has  
compute capability of at least 7.0. Your GPU: NVIDIA GeForce RTX 4050 Laptop  
GPU, compute capability 8.9
```

```
[56]: def inference_grid(img_paths, grid_cols=3):  
    class_names = ['buildings', 'forest', 'glacier', 'mountain', 'sea',  
        ↪ 'street']  
  
    num_images = len(img_paths)  
    grid_rows = math.ceil(num_images / grid_cols) # Calculate number of rows  
    ↪ needed  
  
    plt.figure(figsize=(grid_cols * 4, grid_rows * 4)) # Adjust figure size  
    ↪ dynamically  
  
    for i, img_path in enumerate(img_paths):  
        img = cv2.imread(img_path)  
  
        # Preprocess image  
        img_expanded = np.expand_dims(img, axis=0)  
  
        prediction = model.predict(img_expanded)  
        predicted_class = np.argmax(prediction, axis=1)[0]  
        predicted_label = class_names[predicted_class]  
  
        # Plot image in a grid  
        plt.subplot(grid_rows, grid_cols, i + 1)  
        plt.imshow(img)  
        plt.axis('off')  
        plt.title(f"{predicted_label}", fontsize=12)  
  
    plt.tight_layout() # Adjust layout for better spacing  
    plt.show()
```

```
[60]: parent_dir = "../Dataset/Intel/seg_pred/"  
  
num_images = 9 # Number of images to predict  
  
img_files = random.sample(os.listdir(parent_dir), num_images)  
img_paths = [os.path.join(parent_dir, img) for img in img_files]  
  
inference_grid(img_paths, grid_cols = 3)
```

1/1 [=====] - 0s 17ms/step
 1/1 [=====] - 0s 16ms/step
 1/1 [=====] - 0s 18ms/step
 1/1 [=====] - 0s 20ms/step
 1/1 [=====] - 0s 18ms/step
 1/1 [=====] - 0s 19ms/step
 1/1 [=====] - 0s 18ms/step
 1/1 [=====] - 0s 17ms/step
 1/1 [=====] - 0s 16ms/step


street




street




street



glacier



sea




mountain



buildings



forest



mountain

